



1. Quantization

Full Precision




Low Precision




Computational Fidelity

2. Pruning

Dense Model



Sparse Model



Silent Functions

3. Low-Rank Reduction

Full Rank



$$(d \times d)$$



Low Rank



$$(d \times r, r \times d)$$

Dominant modes

4. Parameter Sharing

Layer 1




Layer 2



⋮

⋮

Layer ℓ



**Sharing
Core**

**Reusable & Layer-
specific Functions**